

# Homework BI-PST.

Andrey Babushkin

**Important note.** The code here is not complete, it only shows the most important parts. A complete script is attached with this document.

---

## Prerequisites

Let's calculate K, L and M and choose an appropriate dataset.

K = 28 (28.05.1996)

L = 9 (Babushkin)

M =  $47 * 28 * 9 \pmod{12} = 47 * 7 * 3 * 12 \pmod{12} = 0$

Dataset: **case0101**

---

## Tasks

- 1) First, we read the dataset, split it and convert Extrinsic and Intrinsic datasets into vectors.

```
1. ds = case0101
2. ds.split = split(ds, ds$Treatment)
3. ds.intrinsic = ds.split$Intrinsic$Score
4. ds.extrinsic = ds.split$Extrinsic$Score
```

Now we print the summary out about both vectors.



```
> summary(ds.intrinsic)
  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
 12.00  17.43   20.40   19.88  22.30   29.70
> var(ds.intrinsic)
[1] 19.70928

> summary(ds.extrinsic)
  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
  5.00  12.15   17.20   15.74  18.95   24.00
> var(ds.extrinsic)
[1] 27.58976
```



So, we have found min, max, mean, median, and variance of both vectors.

2) Now let have a look at densities and empirical distribution functions.

```
> hist(ds.intrinsic, freq=FALSE, breaks = 6)
> plot(ecdf(ds.intrinsic))
```

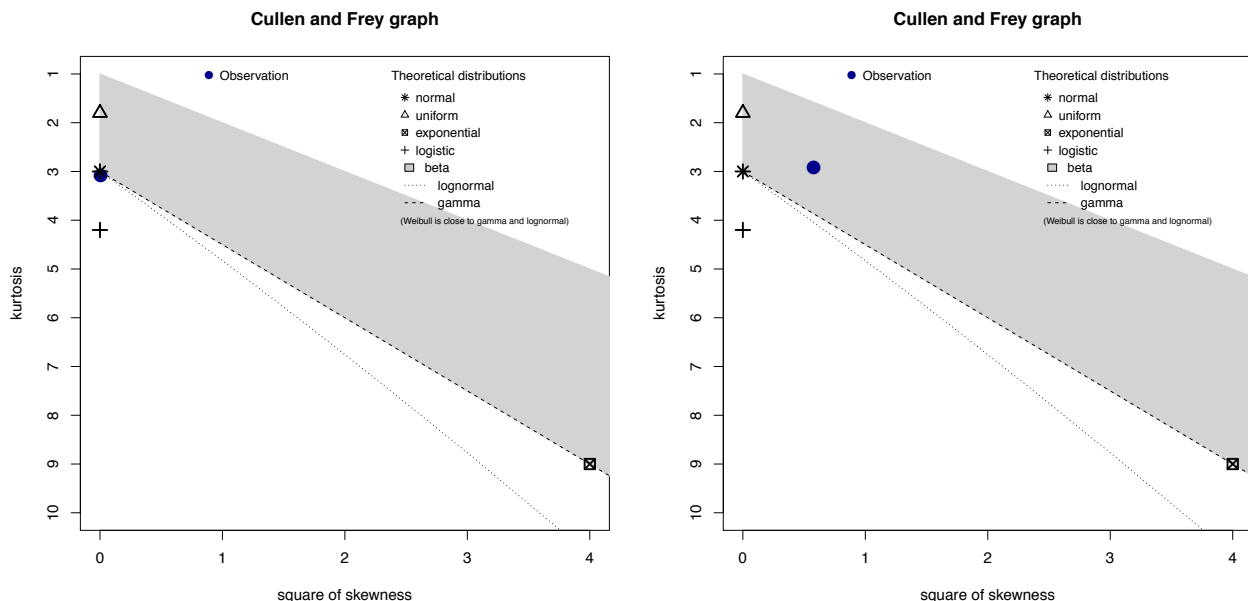


```
> hist(ds.extrinsic, freq=FALSE, breaks = 6)
> plot(ecdf(ds.extrinsic))
```



3) We can see that the distributions looks much like a Gaussian. But let's check our hypothesis using the `descdist` function from the `fitdistrplus` package.

```
> descdist(ds.intrinsic)
> descdist(ds.extrinsic)
```



From these plots we can clearly see that the kurtosis and skewness of our distributions are almost the same as for the Gaussian distribution. Nevertheless, let's fit 3 known distributions (normal, exponential and uniform) and see if they approximate our datasets well.

```
> fit.intrinsic.norm = fitdist(ds.intrinsic, "norm")
> fit.intrinsic.exp = fitdist(ds.intrinsic, "exp")
> fit.intrinsic.unif = fitdist(ds.intrinsic, "unif")

> plot(fit.intrinsic.norm)
> plot(fit.intrinsic.exp)
> plot(fit.intrinsic.unif)

> mu.intrinsic = fit.intrinsic.norm$estimate['mean']
> sigma.intrinsic = fit.intrinsic.norm$estimate['sd']

> x.intrinsic.min = min(ds.intrinsic)
> x.intrinsic.max = max(ds.intrinsic)
> x.intrinsic = seq(x.intrinsic.min, x.intrinsic.max, length =
1000)

> y.intrinsic.norm = dnorm(x.intrinsic,
                           mean = mu.intrinsic,
                           sd = sigma.intrinsic)
> y.intrinsic.exp = dexp(x.intrinsic,
                          rate=fit.intrinsic.exp$estimate['rate'])
```

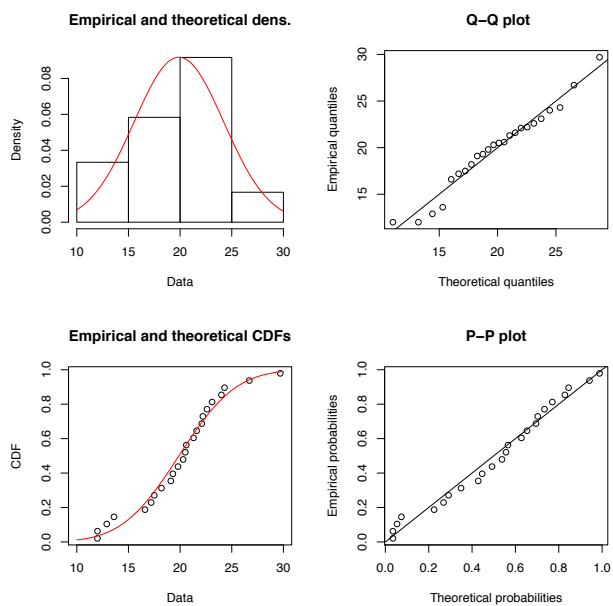
```

> y.intrinsic.unif = dunif(x.intrinsic,
                           min=fit.intrinsic.unif$estimate['min'],
                           max = fit.intrinsic.unif$estimate['max'])

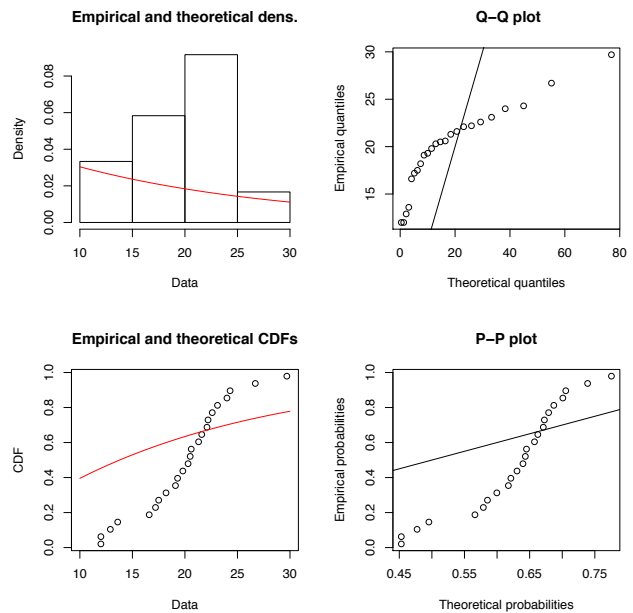
> hist(ds.intrinsic, breaks = 6, freq = FALSE)
> lines(x = x.intrinsic, y = y.intrinsic.norm, col = 'red')
> lines(x = x.intrinsic, y = y.intrinsic.exp, col = 'blue')
> lines(x = x.intrinsic, y = y.intrinsic.unif, col = 'green')

```

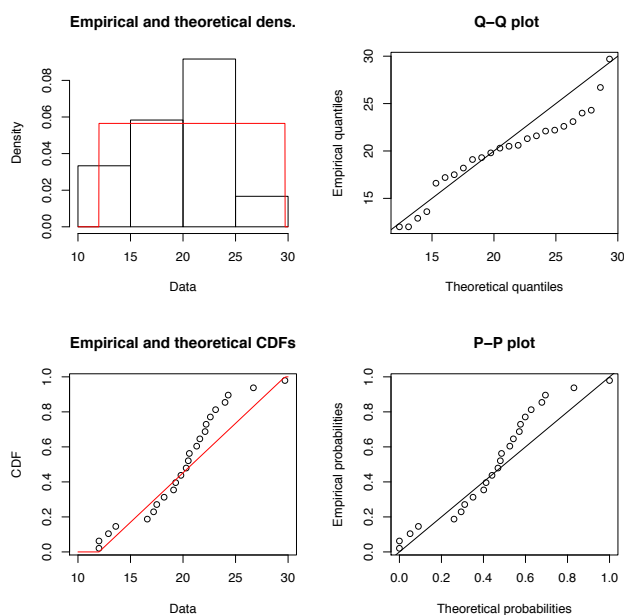
## Normal Intrinsic



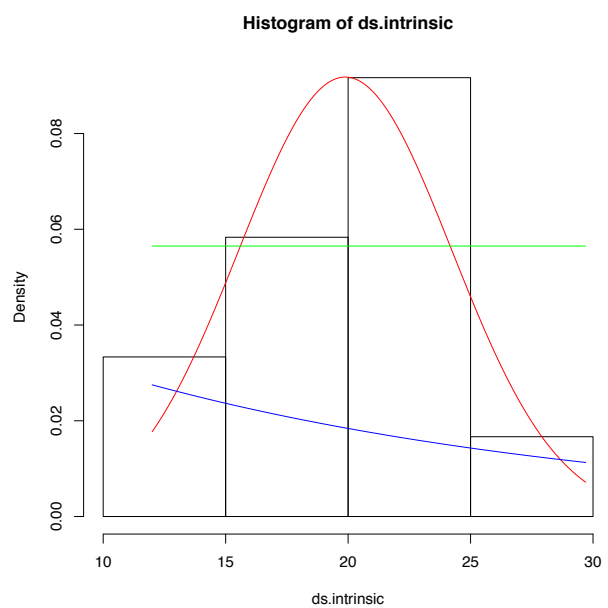
## Exponential Intrinsic



## Uniform Intrinsic

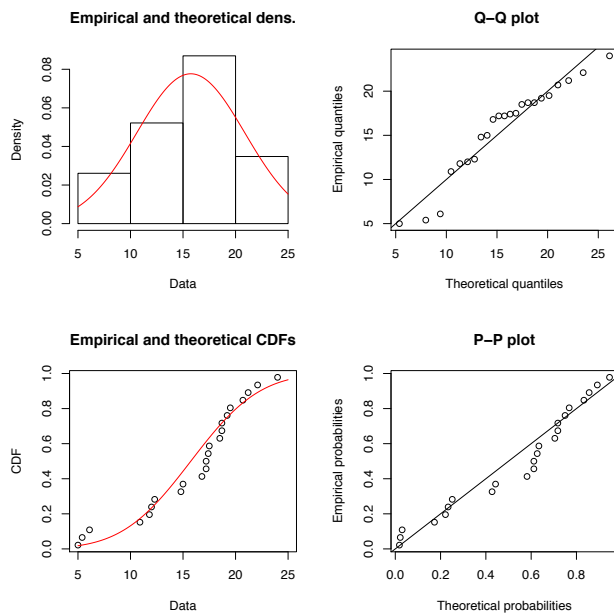


## All fitted distributions

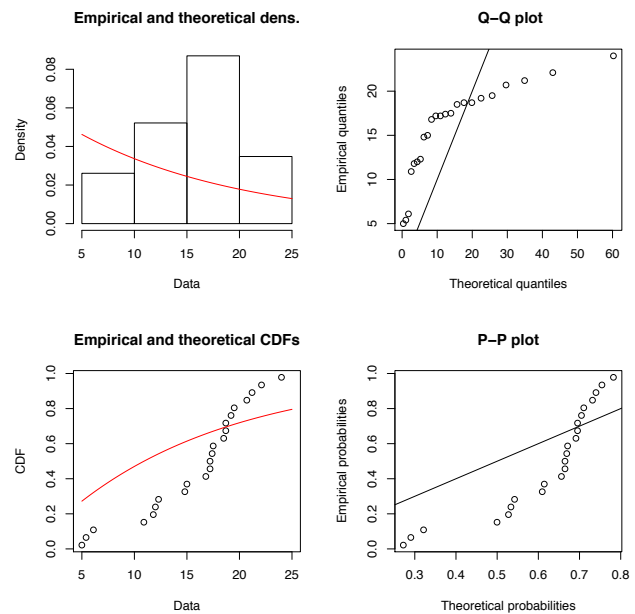


We will do absolutely the same for the Extrinsic dataset with the identical code.

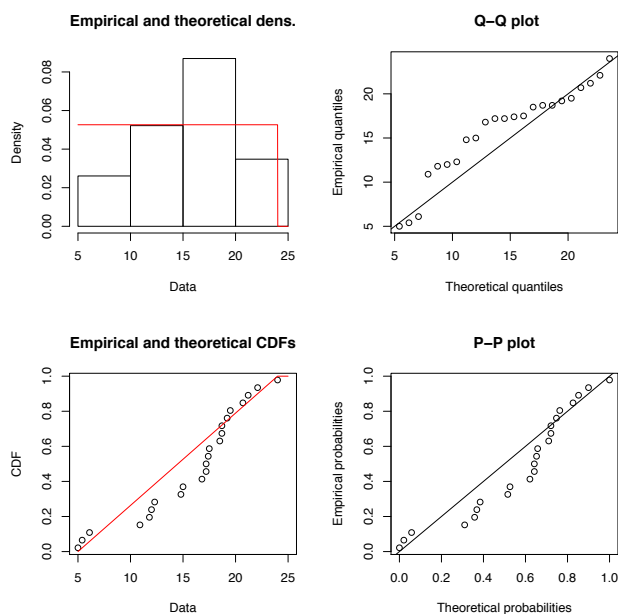
## Normal Extrinsic




## Exponential Extrinsic



## Uniform Extrinsic



## All fitted distributions



From these plot one can clearly see that for both our datasets the best approximation will be a Gaussian distribution.


4) Now we can generate 100 random deviates from the fitted normal distribution and compare histograms of a generated dataset and the actual.

```
> sample.intrinsic.norm = rnorm(100,
                                mean = mu.intrinsic,
                                sd = sigma.intrinsic)

> par(mfrow = c(1, 2))
> hist(ds.intrinsic, breaks = 6,
        freq = FALSE,
        main = "Intrinsic dataset")
> hist(sample.intrinsic.norm, breaks = 6,
        freq = FALSE,
        main = "Generated dataset")

> sample.extrinsic.norm = rnorm(100,
                                mean = mu.extrinsic,
                                sd = sigma.extrinsic)

> par(mfrow = c(1, 2))
> hist(ds.extrinsic, breaks = 6,
        freq = FALSE,
        main = "extrinsic dataset")
> hist(sample.extrinsic.norm, breaks = 6,
        freq = FALSE,
        main = "Generated dataset")
```



So, from these histograms it is seen that the fitted normal distributions approximate our samples quiet well.

5) To calculate confidence intervals for both our distributions we'll use a built-in function in R. Actually, we'll use a side effect of this function - calculating a confidence interval.

```
> siglevel = 0.05

> (conf.int.intrinsic = t.test(ds.intrinsic,
                               conf.level = 1-siglevel)$conf.int)
[1] 18.00869 21.75798
attr("conf.level")
[1] 0.95

> (conf.int.extrinsic = t.test(ds.extrinsic,
                               conf.level = 1-siglevel)$conf.int)
[1] 13.46774 18.01052
attr("conf.level")
[1] 0.95
```

So, our 95% confidence interval for the **Intrinsic** dataset is  
**(18.00869, 21.75798),**

for the **Extrinsic** dataset is  
**(13.46774, 18.01052).**

6) Let's check if the mean value of the datasets is equals to the value of K ( $K = 28$ ). We'll use the 5% significance level interval.

```
> t.test(ds.intrinsic, mu=28)

    One Sample t-test

data:  ds.intrinsic
t = -8.9567, df = 23, p-value = 5.857e-09
alternative hypothesis: true mean is not equal to 28
95 percent confidence interval:
 18.00869 21.75798
sample estimates:
mean of x
 19.88333

> t.test(ds.extrinsic, mu=28)

    One Sample t-test

data:  ds.extrinsic
t = -11.195, df = 22, p-value = 1.491e-10
alternative hypothesis: true mean is not equal to 28
95 percent confidence interval:
 13.46774 18.01052
sample estimates:
mean of x
 15.73913
```

From the output we see that the p-value for the Intrinsic dataset is  $5.857e-09$ , for the Extrinsic dataset is  $1.491e-10$ , which is less than 0.05. That means that we can reject the null-hypothesis, that is, the mean value of the extrinsic data set is not equal to the value of K,  $K = 28$ .



7) The last part of this work will be comparing two datasets or more exactly we will check if the mean values of both dataset are equal on the 95% confidence interval. That will be our null hypothesis (or  $H_0$ ). An alternative hypothesis will be that mean values are not equal.

First of all, before using the two sample t-test, we should evaluate a Fisher's F-test and verify a homogeneity of variances:

```
> var.test(ds.intrinsic, ds.extrinsic)

      F test to compare two variances

data:  ds.intrinsic and ds.extrinsic
F = 0.71437, num df = 23, denom df = 22, p-value = 0.4289
alternative hypothesis: true ratio of variances is not equal to 1
95 percent confidence interval:
 0.3047427 1.6612045
sample estimates:
ratio of variances
      0.7143691
```

We see that the p-value of 0.4289 is greater than 0.05, thus, variances are homogeneous.

So, the variances are homogeneous, our datasets follow the Gaussian distribution, therefore, finally, we can perform a t-test.

```
> t.test(ds.intrinsic, ds.extrinsic, var.equal=TRUE)

      Two Sample t-test

data:  ds.intrinsic and ds.extrinsic
t = 2.9259, df = 45, p-value = 0.005366
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 1.291432 6.996973
sample estimates:
mean of x mean of y
 19.88333  15.73913
```

We have the p-value of 0.005366, which is less than 0.05. So, the  $H_0$  hypothesis can be rejected, mean values of our distributions are different.

BONUS) Let's now draw, just for fun, our 95% confidence intervals.

```
> plot(c(conf.int.intrinsic[1],
        conf.int.intrinsic[2]),
       c(1, 1),
       type = 'l',
       lty = 'solid',
       xlab = 'Confidence intervals',
       ylab = '',
       xlim=c(12, 23),
       ylim=c(0.5, 2),
       col = 'blue')

> points(x = mu.intrinsic, y = 1, pch = 5, col = 'blue')
> segments(x0 = conf.int.intrinsic[1],
          x1 = conf.int.intrinsic[1],
          y0 = 0,
          y1 = 1,
          lty = 'dashed')
> segments(x0 = conf.int.intrinsic[2],
          x1 = conf.int.intrinsic[2],
          y0 = 0,
          y1 = 1,
          lty = 'dashed')
> segments(x0 = conf.int.extrinsic[1],
          y0 = 1.5,
          x1 = conf.int.extrinsic[2],
          y1 = 1.5,
          col = 'red')
> points(x = mu.extrinsic, y = 1.5, pch = 5, col = 'red')
> segments(x0 = conf.int.extrinsic[1],
          x1 = conf.int.extrinsic[1],
          y0 = 0,
          y1 = 1.5,
          lty = 'dashed')
> segments(x0 = conf.int.extrinsic[2],
          x1 = conf.int.extrinsic[2],
          y0 = 0,
          y1 = 1.5,
          lty = 'dashed')

# At the end we'll print two our fitted distributions on one plot.

> plot(x = x.intrinsic,
       y = y.intrinsic.norm,
       col = 'blue', type = 'l', lty = 'solid', xlim=c(5, 30))
> lines(x = x.extrinsic, y = y.extrinsic.norm, col = 'red')
```

